

Experiment No. 3

Aim: Perform Data Modeling.

Problem Statement:

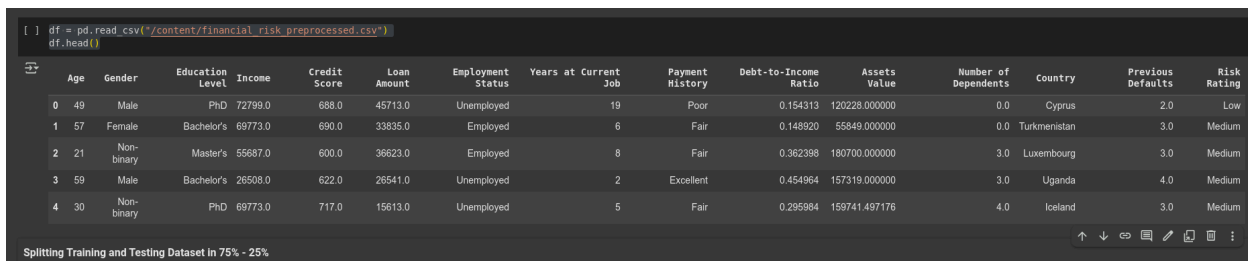
- Partition the data set, for example 75% of the records are included in the training data set and 25% are included in the test data set.
- Use a bar graph and other relevant graph to confirm your proportions.
- Identify the total number of records in the training data set.
- Validate partition by performing a two-sample Z-test.

1. Importing required libraries:

```
[ ] import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
from sklearn.model_selection import train_test_split
from scipy.stats import norm
```

2. Overview of Dataset:

```
df = pd.read_csv("/content/financial_risk_preprocessed.csv")
df.head()
```

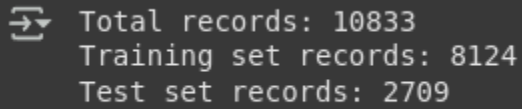


	Age	Gender	Education Level	Income	Credit Score	Loan Amount	Employment Status	Years at Current Job	Payment History	Debt-to-Income Ratio	Assets Value	Number of Dependents	Country	Previous Defaults	Risk Rating
0	49	Male	PhD	72799.0	688.0	45713.0	Unemployed	19	Poor	0.154313	120228.000000	0.0	Cyprus	2.0	Low
1	57	Female	Bachelor's	69773.0	690.0	33835.0	Employed	6	Fair	0.148920	55849.000000	0.0	Turkmenistan	3.0	Medium
2	21	Non-binary	Master's	55687.0	600.0	36623.0	Employed	8	Fair	0.362398	180700.000000	3.0	Luxembourg	3.0	Medium
3	59	Male	Bachelor's	26508.0	622.0	26541.0	Unemployed	2	Excellent	0.454964	157319.000000	3.0	Uganda	4.0	Medium
4	30	Non-binary	PhD	69773.0	717.0	15613.0	Unemployed	5	Fair	0.295984	159741.497176	4.0	Iceland	3.0	Medium

The Financial Risk Assessment Dataset provides detailed information on individual financial profiles. It includes demographic, financial, and behavioral data to assess financial risk. The dataset features various columns such as income, credit score, and risk rating, with intentional imbalances and missing values to simulate real-world scenarios.

3. Splitting Training and Testing Dataset in 75% - 25% :

```
train, test = train_test_split(df, test_size=0.25, random_state=42)
print(f'Total records: {len(df)}')
print(f'Training set records: {len(train)}')
print(f'Test set records: {len(test)}')
```

A terminal window with a dark background and light gray text. It shows the output of the print statements from the code above: 'Total records: 10833', 'Training set records: 8124', and 'Test set records: 2709'.



```
↵ Total records: 10833
   Training set records: 8124
   Test set records: 2709
```

The dataset is divided into 75% for training and 25% for testing.

4. Plotting graph of Training and Testing Dataset

```
plt.figure(figsize=(8, 5))
plt.bar(["Training Set", "Test Set"], [len(train), len(test)], color=['blue', 'orange'])
plt.ylabel("Number of Records")
plt.title("Dataset Partitioning")
plt.show()
```

```
plt.figure(figsize=(6, 6))
plt.pie([len(train), len(test)], labels=["Training Set", "Test Set"], autopct="%1.1f%%",
        colors=['blue', 'orange'])
plt.title("Training vs. Test Data Distribution")
plt.show()
```



From above graph we can see that our data is properly partitioned into 75% training data and 25% testing data

5. Performing Z-Test

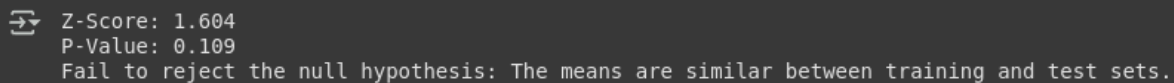
```
column_name = df.columns[0]
train_mean = train[column_name].mean()
test_mean = test[column_name].mean()
train_std = train[column_name].std()
test_std = test[column_name].std()
n_train = len(train)
n_test = len(test)

# Compute Z-score
z_score = (train_mean - test_mean) / np.sqrt((train_std**2 / n_train) + (test_std**2 / n_test))

# Compute p-value
p_value = 2 * (1 - norm.cdf(abs(z_score)))

print(f'Z-Score: {z_score:.3f}')
print(f'P-Value: {p_value:.3f}')

# Interpretation
alpha = 0.05
if p_value < alpha:
    print("Reject the null hypothesis: The means of the two groups are significantly different.")
else:
    print("Fail to reject the null hypothesis: The means are similar between training and test sets.")
```



```
Z-Score: 1.604
P-Value: 0.109
Fail to reject the null hypothesis: The means are similar between training and test sets.
```

Performing two sample z-test on 'Age' columns. Given that the null hypothesis was not rejected, the data split is statistically valid

6. Performing Correlation Test

```
# Calculate Pearson correlation for all numerical columns
correlation_matrix = train.corr(numeric_only=True)
print("Correlation Matrix (Training Set):\n", correlation_matrix)
```

```
# Visualize the correlation matrix
```

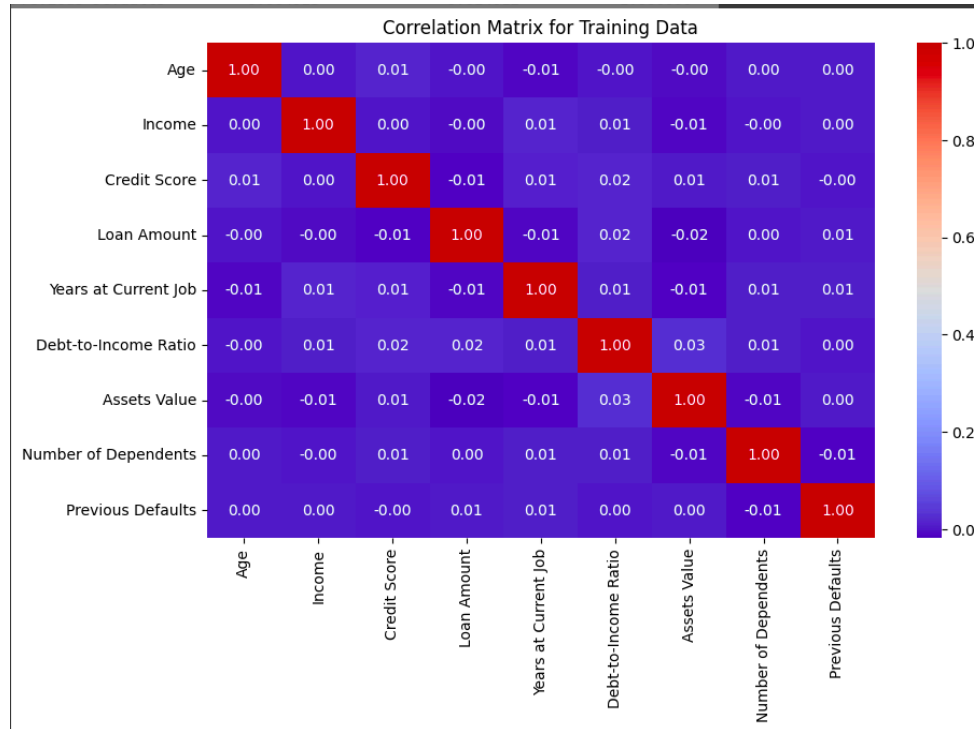
```
import seaborn as sns
import matplotlib.pyplot as plt

plt.figure(figsize=(10, 6))
sns.heatmap(correlation_matrix, annot=True, cmap="coolwarm", fmt=".2f")
plt.title("Correlation Matrix for Training Data")
plt.show()
```

Correlation Matrix (Training Set):

	Age	Income	Credit Score	Loan Amount	\
Age	1.000000	0.001015	0.013612	-0.000546	
Income	0.001015	1.000000	0.001392	-0.004547	
Credit Score	0.013612	0.001392	1.000000	-0.013475	
Loan Amount	-0.000546	-0.004547	-0.013475	1.000000	
Years at Current Job	-0.006949	0.013522	0.011674	-0.007857	
Debt-to-Income Ratio	-0.001505	0.007965	0.015618	0.016472	
Assets Value	-0.003270	-0.006882	0.005112	-0.018345	
Number of Dependents	0.004430	-0.000333	0.007754	0.001100	
Previous Defaults	0.001992	0.001911	-0.001708	0.005074	
	Years at Current Job	Debt-to-Income Ratio	\		
Age	-0.006949	-0.001505			
Income	0.013522	0.007965			
Credit Score	0.011674	0.015618			
Loan Amount	-0.007857	0.016472			
Years at Current Job	1.000000	0.008937			
Debt-to-Income Ratio	0.008937	1.000000			
Assets Value	-0.012742	0.026238			
Number of Dependents	0.008165	0.007587			
Previous Defaults	0.007654	0.000719			
	Assets Value	Number of Dependents	Previous Defaults		
Age	-0.003270	0.004430	0.001992		
Income	-0.006882	-0.000333	0.001911		
Credit Score	0.005112	0.007754	-0.001708		
Loan Amount	-0.018345	0.001100	0.005074		
Years at Current Job	-0.012742	0.008165	0.007654		
Debt-to-Income Ratio	0.026238	0.007587	0.000719		
Assets Value	1.000000	-0.006907	0.004923		
Number of Dependents	-0.006907	1.000000	-0.014012		
Previous Defaults	0.004923	-0.014012	1.000000		

The negligible correlation values indicate an absence of a meaningful relationship between the columns, a finding that is further supported by the correlation heatmap.



7. Performing Chi-Squared Test:

Create a contingency table for Education Level and Employment Status

```
contingency_table = pd.crosstab(train['Education Level'], train['Employment Status'])
```

```
print("Contingency Table:\n", contingency_table)
```

```
print("\n\n")
```

```
from scipy.stats import chi2_contingency
```

Perform the chi-squared test

```
chi2, p, dof, expected = chi2_contingency(contingency_table)
```

Display the results

```
print(f'Chi-Squared Statistic: {chi2:.3f}')
```

```
print(f'p-value: {p:.3f}')
```

```
print(f'Degrees of Freedom: {dof}')
```

```
print("Expected Frequencies:\n", pd.DataFrame(expected, index=contingency_table.index,
columns=contingency_table.columns))
```

```
print("\n\n")
```

Interpret the p-value

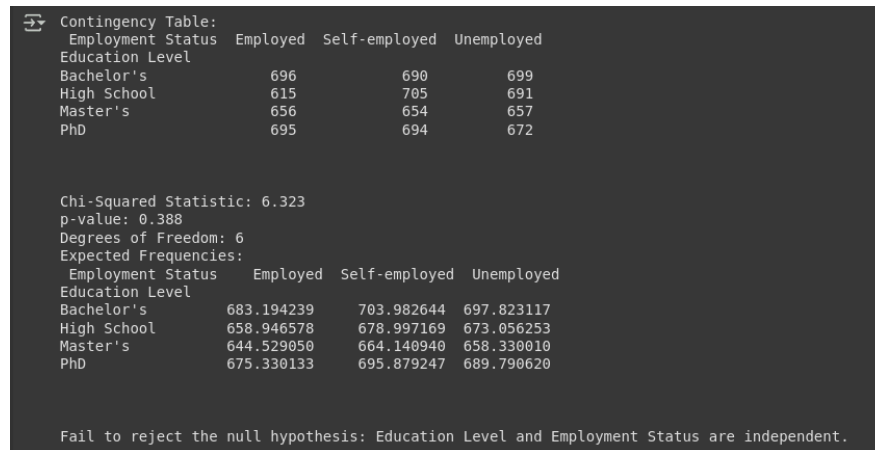
```
alpha = 0.05
```

```
if p < alpha:
```

```
    print("Reject the null hypothesis: Education Level and Employment Status are dependent.")
```

else:

```
print("Fail to reject the null hypothesis: Education Level and Employment Status are independent.")
```



The image shows a Jupyter Notebook output. It contains a contingency table for 'Education Level' (rows) and 'Employment Status' (columns: Employed, Self-employed, Unemployed). Below the table, it displays the results of a Chi-Squared test: Chi-Squared Statistic: 6.323, p-value: 0.388, Degrees of Freedom: 6. It also shows the expected frequencies for each cell. At the bottom, a message states: 'Fail to reject the null hypothesis: Education Level and Employment Status are independent.'

Employment Status	Employed	Self-employed	Unemployed
Education Level			
Bachelor's	696	690	699
High School	615	705	691
Master's	656	654	657
PhD	695	694	672

Employment Status	Employed	Self-employed	Unemployed
Education Level			
Bachelor's	683.194239	703.982644	697.823117
High School	658.946578	678.997169	673.056253
Master's	644.529050	664.140940	658.330010
PhD	675.330133	695.879247	689.790620

Chi-Squared Statistic: 6.323
p-value: 0.388
Degrees of Freedom: 6
Expected Frequencies:

Fail to reject the null hypothesis: Education Level and Employment Status are independent.

The Chi-Squared test was performed on columns 'Education Level' and 'Employment Status'. Since the test results do not reject the null hypothesis, it can be concluded that 'Education Level' and 'Employment Status' are independent variables according to the dataset.

8. Download partitioned Training and Testing dataset

```
train.to_csv("financial_risk_train_data.csv", index=False)
test.to_csv("financial_risk_test_data.csv", index=False)
from google.colab import files
files.download("financial_risk_train_data.csv")
files.download("financial_risk_test_data.csv")
```

We can use partitioned dataset for training and testing purposes

Conclusion:

We loaded the data into a Colab notebook and split it into 75% for training and 25% for testing. To verify the partition, we plotted a bar graph and a pie chart, which confirmed that the data was split correctly. Next, we performed a Z-test to assess the validity of the partition, and the results showed that the partition was valid, as the null hypothesis was not rejected. We then examined the correlation between all columns using a correlation heatmap, which revealed no significant correlation between the columns. Finally, we conducted a chi-square test on the 'Education Level' and 'Employment Status' columns, and the results indicated that the two features are independent, as the null hypothesis was not rejected.